



Machine Learning Engineer Challenge

Thanks a lot for your interest in the Machine Learning Engineer position in our team! We are very excited to get to know you better and introduce you to our work.

Our team

At Vocads, we are a team of builders at heart: our motto is to be pragmatic and always go all the way to production. We strive to read the latest research papers, adapt them to our specific problems, build prototypes, and finally craft fully packaged products off of them.

Context

Our team is currently building an ensemble of natural language processing models to power Vocads conversational AI platform.

The current prediction pipeline is partly based on a recent research paper in the field of NLP: sentenceBERT <https://arxiv.org/pdf/1908.10084.pdf>.

We have built this challenge to give you a taste of the kind of work we do in our team. Don't be worried if your prediction performance is not great, it is expected in such a limited amount of time.

If you feel like there is something unclear, please reach out to us with your questions. There are no bad points for doing so, quite the contrary. We hope you will have fun working on this task, just like we do every day!

The challenge

NB: you do not need extensive knowledge in NLP for this test. No one can be an expert in all the subfields of Machine Learning. However, we expect you to cope on your own by reading the documentation found in the researcher's repository on Github. We strongly advise you to use Google Collab to get access to free GPUs, unless you have one at your disposal.

Based on this repository <https://github.com/UKPLab/sentence-transformers> and the 20 Newsgroups dataset from sklearn

https://scikit-learn.org/0.19/datasets/twenty_newsgroups.html, you will build a sentence classifier in Python:

1. In a Notebook, create a custom loading class, compatible with Pytorch DataLoader, that generates training triplets (anchor, positive example, negative example) from 20 Newsgroups. You might want to take a look at the SentenceLabelDataset class
https://github.com/UKPLab/sentence-transformers/blob/6fcdfb30f9dfcc5fb978c97ce02941a7aa6ba63/sentence_transformers/datasets/SentenceLabelDataset.py.

You should come up with a strategy to generate triplets that will be the most helpful / insightful for the model to train with.

2. Build a training pipeline and fine-tune a distilbert-base-nli-mean-tokens model with your custom loading class, using the TripletLoss loss function. Since fine-tuning is quite time-consuming, even on a GPU, you can go for a single epoch. Your triplet generation strategy is what matters to us.
3. Do some research online to find an Approximate Nearest Neighbor library. There are many of them and there really is no wrong choice. But we expect you to explain your pick in a few words.
4. Build a basic prediction pipeline:
 - a. Vectorize the training set with your fine-tuned sBERT model
 - b. Index all these vectors with your ANN library
 - c. Build a barebone kNN classifier where new text input gets predicted the same label as that of the closest neighbor from the index
 - d. Benchmark this pipeline with the test set
 - e. Compare the results with the pretrained sBERT model
5. Create a simple Python REST API that serves this prediction via a “/predict” route.

Given some input text, it should return one of the 20 Newsgroups labels as a prediction.

We expect your code to be wrapped in a docker container via a Dockerfile.

Please send us back:

- your code for the API in a private Github/Gitlab repo
- a link to your Colab Notebook (or add your jupyter notebook to the repo)
- a pdf version of your notebook (just in case)